

Distributed Programming

Problems & Solutions

Agenda

- Concurrency Control
 - Software Algorithms
 - Semaphores (OS Support)
 - Monitors (Compiler Support)
 - Distributed Mutual Exclusion Algorithms
- Distributed Coordination
 - Synchronized Clocks
 - Logical Clocks
 - Election Algorithms

Strict Alternation

```
while (TRUE) {  
    while (turn != 0)      /* loop */ ;  
    critical_region( );  
    turn = 1;  
    noncritical_region( );  
}
```

(a)

```
while (TRUE) {  
    while (turn != 1)      /* loop */ ;  
    critical_region( );  
    turn = 0;  
    noncritical_region( );  
}
```

(b)

First me, then
you

Problems with strict alternation

- Employs busy waiting-while waiting for the cr, a process spins
- If one process is outside the cr and it is its turn, then other process has to wait until outside guy finishes both outside AND inside (cr) work

Peterson's Solution (1981)

```
#define FALSE 0
#define TRUE 1
#define N      2                /* number of processes */

int turn;                       /* whose turn is it? */
int interested[N];              /* all values initially 0 (FALSE) */

void enter_region(int process);  /* process is 0 or 1 */
{
    int other;                  /* number of the other process */

    other = 1 - process;        /* the opposite of process */
    interested[process] = TRUE; /* show that you are interested */
    turn = process;             /* set flag */
    while (turn == process && interested[other] == TRUE) /* null statement */ ;
}

void leave_region(int process)   /* process: who is leaving */
{
    interested[process] = FALSE; /* indicate departure from critical region */
}
```

Peterson

- Process 0 & 1 try to get in simultaneously
- Last one in sets turn: say it is process 1
- Process 0 enters (turn = process is False)

Semaphores

- Semaphore is an integer variable
- Used to sleeping processes/wakeups
- Two operations, **down/P/Wait** and **up/V/Signal**
- Down checks semaphore. If not zero, decrements semaphore. If zero, process goes to sleep
- Up increments semaphore. If more than one process asleep, one is chosen randomly and enters critical region (first does a down)
- **ATOMIC IMPLEMENTATION**-interrupts disabled

Producer Consumer with Semaphores

- 3 semaphores: full, empty and mutex
- Full counts full slots (initially 0)
- Empty counts empty slots (initially N)
- Mutex protects variable which contains the items produced and consumed

Producer Consumer with semaphores

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;

void producer(void)
{
    int item;

    while (TRUE) {
        item = produce_item();
        down(&empty);
        down(&mutex);
        insert_item(item);
        up(&mutex);
        up(&full);
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {
        down(&full);
        down(&mutex);
        item = remove_item();
        up(&mutex);
        up(&empty);
        consume_item(item);
    }
}
. . . }
```

/* number of slots in the buffer */
/* semaphores are a special kind of int */
/* controls access to critical region */
/* counts empty buffer slots */
/* counts full buffer slots */

/* TRUE is the constant 1 */
/* generate something to put in buffer */
/* decrement empty count */
/* enter critical region */
/* put new item in buffer */
/* leave critical region */
/* increment count of full slots */

/* infinite loop */
/* decrement full count */
/* enter critical region */
/* take item from buffer */
/* leave critical region */
/* increment count of empty slots */
/* do something with the item */

Monitors

- Easy to make a mess of things using mutexes and condition variables. Little errors cause disasters.
- Producer consumer with semaphores-
interchange two downs in producer code causes deadlock
- Monitor is a **language construct** which enforces mutual exclusion and blocking mechanism
- C does not have monitor

Monitors

- Monitor consists of {procedures, data structures, and variables} grouped together in a “module”
- A process can call procedures inside the monitor, but cannot directly access the stuff inside the monitor

Monitor-a picture

```
monitor example
  integer i;
  condition c;

  procedure producer( );
  .
  .
  end;

  procedure consumer( );
  .      .      .
  end;
end monitor;
```

Onwards

- In a monitor it is the job of the compiler, not the programmer to enforce mutual exclusion.
- Only one process at a time can be in the monitor
When a process calls a monitor, the first thing done is to check if another process is in the monitor. If so, calling process is suspended.
- Need to enforce blocking as well –
 - use condition variables
 - Use wait , signal ops on cv's

Condition Variables

- Monitor discovers that it can't continue (e.g. buffer is full), issues a signal on a condition variable (e.g. full) causing process (e.g. producer) to block
- Another process is allowed to enter the monitor (e.g. consumer). This process can issue a signal, causing blocked process (producer) to wake up
- Process issuing signal leaves monitor

Producer Consumer Monitor

```
monitor ProducerConsumer
  condition full, empty;
  integer count;

  procedure insert(item: integer);
  begin
    if count = N then wait(full);
    insert_item(item);
    count := count + 1;
    if count = 1 then signal(empty)
  end;

  function remove: integer;
  begin
    if count = 0 then wait(empty);
    remove = remove_item;
    count := count - 1;
    if count = N - 1 then signal(full)
  end;

  count := 0;
end monitor;

procedure producer;
begin
  while true do
    begin
      item = produce_item;
      ProducerConsumer.insert(item)
    end
  end;

procedure consumer;
begin
  while true do
    begin
      item = ProducerConsumer.remove;
      consume_item(item)
    end
  end;
end;
```

Monitors: Good vs Bad

- The good- No messy direct programmer control of semaphores
- The bad- You need a language which supports monitors (Java - synchronized).
- OS's are written in C

Semaphores: Good vs Bad

- The good- Easy to implement
- The bad- Easy to mess up

Reality

- Monitors and semaphores only work for shared memory
- Don't work for multiple CPU's which have their own private memory, e.g. workstations on an Ethernet

Message Passing

- Used when memory is not shared
- Information exchange between machines
- Two primitives
 - `Send(destination, &message)`
 - `Receive(source, &message)`
- Lots of design issues
 - Message loss
 - acknowledgements, time outs deal with loss
 - Authentication-how does a process know the identity of the sender? For sure, that is

Producer Consumer Using Message Passing

- Consumer sends N empty messages to producer
- Producer fills message with data and sends to consumer

Producer-Consumer Problem with Message Passing (1)

```
#define N 100                                /* number of slots in the buffer */

void producer(void)
{
    int item;
    message m;                                /* message buffer */

    while (TRUE) {
        item = produce_item();                /* generate something to put in buffer */
        receive(consumer, &m);                /* wait for an empty to arrive */
        build_message(&m, item);              /* construct a message to send */
        send(consumer, &m);                   /* send item to consumer */
    }
}
```

Producer-Consumer Problem with Message Passing (2)

• • •

```
void consumer(void)
{
    int item, i;
    message m;

    for (i = 0; i < N; i++) send(producer, &m); /* send N empties */
    while (TRUE) {
        receive(producer, &m);                /* get message containing item */
        item = extract_item(&m);               /* extract item from message */
        send(producer, &m);                    /* send back empty reply */
        consume_item(item);                    /* do something with the item */
    }
}
```

Message Passing Approaches

- Have unique ID for address of recipient process
- Mailbox
 - In producer consumer, have one for the producer and one for the consumer
- No buffering-sending process blocks until the receive happens. Receiver blocks until send occurs (Rendezvous)

Next

Distributed Algorithms